

Linear Regression

Art Tay

Gradient Derivation

a)

$$\begin{aligned}\frac{\partial}{\partial \beta_0} \left[\frac{1}{2m} \sum_{i=0}^m (\beta_0 + \beta_1 x^{(i)} - y^{(i)})^2 \right] &= \frac{\partial}{\partial \beta_0} \left[\frac{1}{2m} \sum_{i=0}^m (\beta_0 + \beta_1 x^{(i)} - y^{(i)})^2 \right] \\ &= \frac{1}{2m} \sum_{i=0}^m \left[2(\beta_0 + \beta_1 x^{(i)} - y^{(i)}) \cdot 1 \right] \\ &= \frac{1}{m} \sum_{i=0}^m \beta_0 + \beta_1 \frac{1}{m} \sum_{i=0}^m x^{(i)} - \frac{1}{m} \sum_{i=0}^m y^{(i)} \\ &= \beta_0 + \beta_1 \bar{x} - \bar{y}\end{aligned}$$

b)

$$\begin{aligned}\frac{\partial}{\partial \beta_1} \left[\frac{1}{2m} \sum_{i=0}^m (\beta_0 + \beta_1 x^{(i)} - y^{(i)})^2 \right] &= \frac{1}{2m} \sum_{i=0}^m \frac{\partial}{\partial \beta_1} \left[(\beta_0 + \beta_1 x^{(i)} - y^{(i)})^2 \right] \\ &= \frac{1}{2m} \sum_{i=0}^m \left[2x^{(i)} (\beta_0 + \beta_1 x^{(i)} - y^{(i)}) \right] \\ &= \frac{1}{m} \sum_{i=0}^m \left[\beta_0 x^{(i)} + \beta_1 x^{(i)2} - x^{(i)} y^{(i)} \right] \\ &= \beta_0 \frac{1}{m} \sum_{i=0}^m x^{(i)} + \beta_1 \frac{1}{m} \sum_{i=0}^m x^{(i)2} - \frac{1}{m} \sum_{i=0}^m x^{(i)} y^{(i)} \\ &= \beta_0 \bar{x} + \beta_1 \bar{x}^2 - \bar{x}\bar{y}\end{aligned}$$

Linear Regression by Gradient Decent

```
# Generates linear data with normal residuals
set.seed(123)
x <- rnorm(n = 30)

epsilon <- rnorm(n = 30)

y <- 5*x + 1 + epsilon
```

```
summary(lm(y ~ x))
```

```
##
## Call:
## lm(formula = y ~ x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6085 -0.5056 -0.2152  0.6932  2.0118
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.1720     0.1534   7.639 2.54e-08 ***
## x             4.8660     0.1589  30.629 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8393 on 28 degrees of freedom
## Multiple R-squared:  0.971, Adjusted R-squared:  0.97
## F-statistic: 938.1 on 1 and 28 DF, p-value: < 2.2e-16
```

```
# Calculates the mean squared error for a simple linear regression model.
# @param x - a vector of the explanatory variable.
# @param y - a vector of the response variable.
# @param beta_0 - the intercept value for the current SLR model.
# @param beta_1 - the slope value for the current SLR model.
# @return - sum total mean squared error (y_hat - y)^2
slr_mse <- function(x, y, beta_0, beta_1){
  cost <- ((beta_1 * x + beta_0) - y)^2
  return(sum(cost))
}
```

```
# Calculates the slope and intercept values for SLR
# or simple linear regression.
# @param x - a vector of the explanatory variable.
# @param y - a vector of the response variable.
# @param alpha - the learning rate.
# @return betas - a vector containing the calculated betas.
slr_gradient_desc <- function(x, y, alpha){

  # Summary statistic calculations.
  # Helps to calculate the gradient faster.
  x_bar <- mean(x)
  y_bar <- mean(y)
  xy_bar <- mean(x*y)
  x_sqbar <- mean(x^2)

  # initial guess for beta_0 and beta_1.
  beta_0 <- y_bar
  beta_1 <- 0

  # A counter to determine is the error is unchanging.
  # This is the Loop-Control-Variable (LCV).
  count_same <- 0

  # Iterate 100 times or until the cost remains unchanged for 10 iterations.
```

```

for(i in 1:1000){

  # Stop the loop if the LCV >= 10.
  if(count_same >= 10){
    break
  }

  # Cost prior to beta adjustment.
  cost_start <- slr_mse(x, y, beta_0, beta_1)

  # Calculate gradient values.
  g_0 <- beta_0 + (beta_1 * x_bar) - y_bar
  g_1 <- (beta_0 * x_bar) + (beta_1 * x_sqbar) - xy_bar

  # Update betas.
  beta_0 <- beta_0 - (alpha * g_0)
  beta_1 <- beta_1 - (alpha * g_1)

  # If the cost is unchanged add 1 to the LCV.
  if(cost_start == slr_mse(x, y, beta_0, beta_1)){
    count_same <- count_same + 1
  }
}

return(c(beta_0 = round(beta_0, 4),
        beta_1 = round(beta_1, 4),
        iterations = i))
}

```

```
slr_gradient_desc(x, y, alpha = 0.1)
```

```
##      beta_0      beta_1 iterations
##      1.172      4.866    256.000
```

Linear Model on Airbnb Data