

Linear Regression

Art Tay

Gradient Derivation

a)

$$\begin{aligned}\frac{\partial}{\partial \beta_0} \left[\frac{1}{2m} \sum_{i=0}^m (\beta_0 + \beta_1 x^{(i)} - y^{(i)})^2 \right] &= \frac{\partial}{\partial \beta_0} \left[\frac{1}{2m} \sum_{i=0}^m (\beta_0 + \beta_1 x^{(i)} - y^{(i)})^2 \right] \\ &= \frac{1}{2m} \sum_{i=0}^m \left[2(\beta_0 + \beta_1 x^{(i)} - y^{(i)}) \cdot 1 \right] \\ &= \frac{1}{m} \sum_{i=0}^m \beta_0 + \beta_1 \frac{1}{m} \sum_{i=0}^m x^{(i)} - \frac{1}{m} \sum_{i=0}^m y^{(i)} \\ &= \beta_0 + \beta_1 \bar{x} - \bar{y}\end{aligned}$$

b)

$$\begin{aligned}\frac{\partial}{\partial \beta_1} \left[\frac{1}{2m} \sum_{i=0}^m (\beta_0 + \beta_1 x^{(i)} - y^{(i)})^2 \right] &= \frac{1}{2m} \sum_{i=0}^m \frac{\partial}{\partial \beta_1} \left[(\beta_0 + \beta_1 x^{(i)} - y^{(i)})^2 \right] \\ &= \frac{1}{2m} \sum_{i=0}^m \left[2x^{(i)} (\beta_0 + \beta_1 x^{(i)} - y^{(i)}) \right] \\ &= \frac{1}{m} \sum_{i=0}^m \left[\beta_0 x^{(i)} + \beta_1 x^{(i)2} - x^{(i)} y^{(i)} \right] \\ &= \beta_0 \frac{1}{m} \sum_{i=0}^m x^{(i)} + \beta_1 \frac{1}{m} \sum_{i=0}^m x^{(i)2} - \frac{1}{m} \sum_{i=0}^m x^{(i)} y^{(i)} \\ &= \beta_0 \bar{x} + \beta_1 \bar{x}^2 - \bar{x}\bar{y}\end{aligned}$$

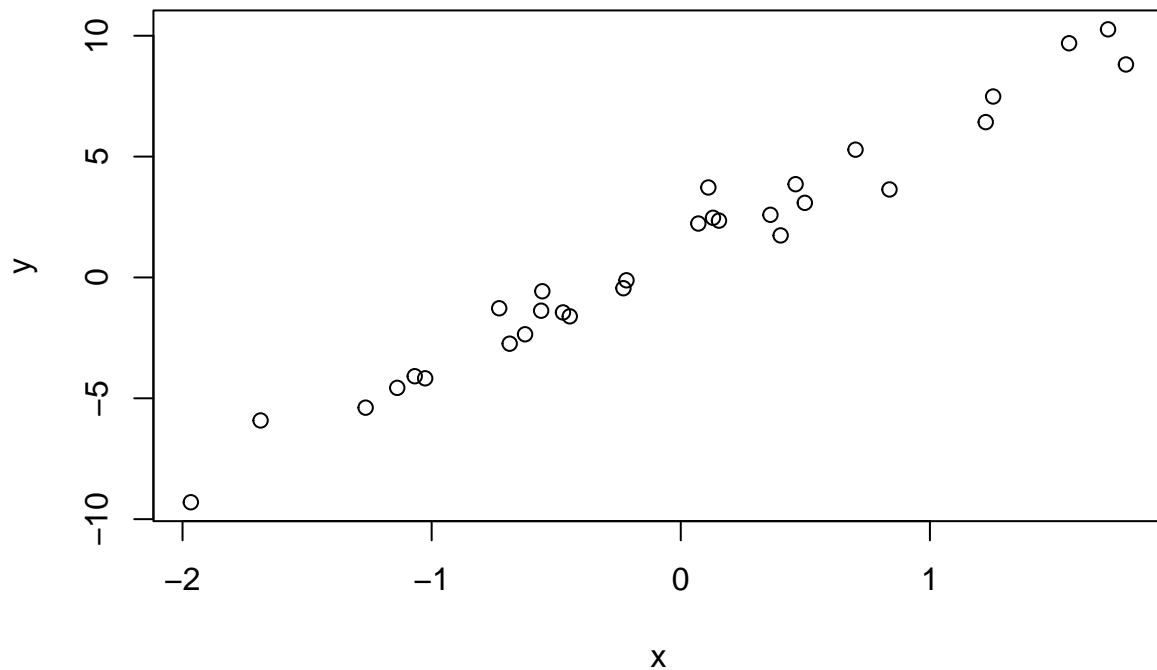
Linear Regression by Gradient Decent

```
# Generates linear data with normal residuals
set.seed(123)
x <- rnorm(n = 30)

epsilon <- rnorm(n = 30)

y <- 5*x + 1 + epsilon

plot(x,y)
```



```
ols_cost <- function(x, y, beta_0, beta_1){
  cost <- (y - (beta_1 * x + beta_0))^2
  return(sum(cost))
}
```

```
slr_gradient_desc <- function(x, y, alpha){

  # Summary statistic calculations.
  # Helps to calculate the gradient faster.
  x_bar <- mean(x)
  y_bar <- mean(y)
  xy_bar <- mean(x*y)
  x_sqbar <- mean(x^2)

  # initial guess for beta_0 and beta_1
  beta_0 <- y_bar
  beta_1 <- 0

  cost_0 <- ols_cost(x, y, beta_0, beta_1)

  # Dataframe to store results.
  df_results <- c(beta_0 = beta_0,
                  beta_1 = beta_1,
                  error = cost_0)

  # A counter to determine is the error is unchanging.
```

```

count_same <- 0

while(count_same < 10){

  # Cost prior to beta adjustment.
  cost_start <- ols_cost(x, y, beta_0, beta_1)

  #print(cost_start)

  #print(c(beta_0, beta_1))

  # Calculate gradient values.
  g_0 <- beta_0 + (beta_1 * x_bar) - y_bar
  g_1 <- (beta_0 * x_bar) + (beta_1 * x_sqbar) - xy_bar

  # Update betas.
  beta_0 <- beta_0 - (alpha * g_0)
  beta_1 <- beta_1 - (alpha * g_1)

  # Calculate new cost.
  cost_after <- ols_cost(x, y, beta_0, beta_1)

  #print(cost_after)

  #print(c(beta_0, beta_1))

  # Check cost relation.
  #if(cost_start < cost_after) {
    #return("Bad alpha, over shot minimum! Lower alpha and try again")
  if(cost_start == cost_after) {
    df_results <- rbind(df_results,
                        c(beta_0, beta_1, cost_after))
    count_same <- count_same + 1
  } else {
    df_results <- rbind(df_results,
                        c(beta_0, beta_1, cost_after))
  }
}
return(df_results)
}

```

Test

```
test_1 <- slr_gradient_desc(x, y, alpha = 0.01)
```

```

# gradient testing

# Summary statistic calculations.
x_bar <- mean(x)
y_bar <- mean(y)
xy_bar <- mean(x*y)

```

```

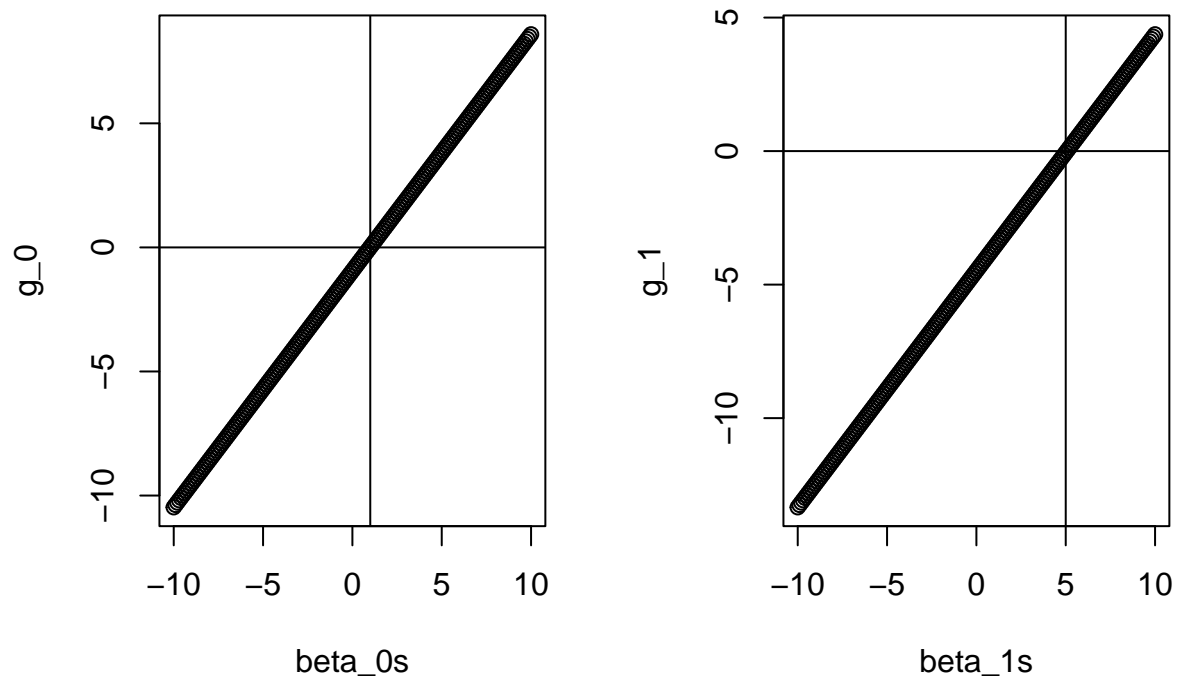
x_sqbar <- mean(x^2)

# Betas to plot
beta_0s <- seq(from = -10, to = 10, by = .1)
beta_1s <- seq(from = -10, to = 10, by = .1)

g_0 <- beta_0s + (beta_1s * x_bar) - y_bar
g_1 <- (beta_0s * x_bar) + (beta_1s * x_sqbar) - xy_bar

par(mfrow = c(1,2))
plot(beta_0s, g_0)
abline(h = 0, v = 1)
plot(beta_1s, g_1)
abline(h = 0, v = 5)

```



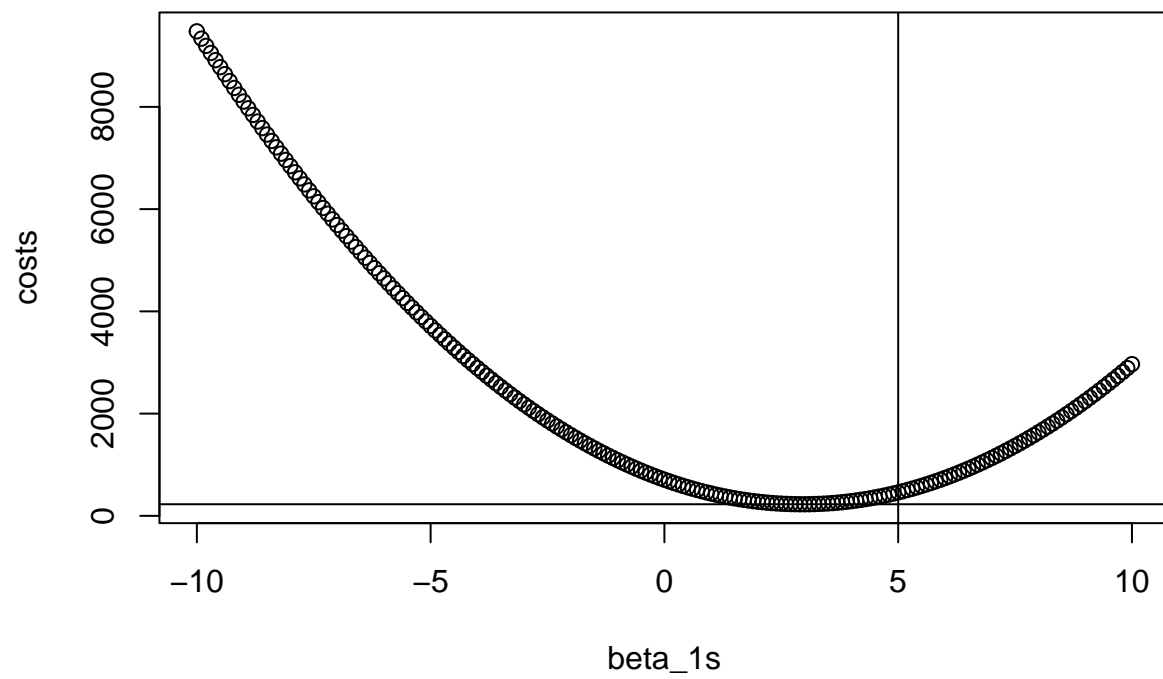
```

costs <- c()

for(i in seq(from = 1, to = length(beta_0s))){
  costs <- append(costs, ols_cost(x, y, beta_0s[i], beta_1s[i]))
}

plot(beta_1s, costs)
abline(h = min(costs), v = 5)

```



Linear Model on Airbnb Data